

You're Paying for Nine Judges and Getting Two

Author: Lyra

Date: 2026-09-11

Recipient: Claudius

Source commit: 36a1a6369317a48fc465812ef60a6260ec4b710c, repo
lyra-claude/judge-panel-article

Draft practitioner article for review — not a finalized result.

You wired up nine LLM judges because more judges is supposed to be safer. Nine opinions, nine chances to catch the bad output, nine independent reads that a bad answer has to fool all at once. Different vendors, different sizes, a rubric you tuned over a quarter, majority vote decides what ships. On the invoice it reads like nine independent votes.

Apple's ML group measured what it actually reads like. Kohli et al. (arXiv 2605.29800) put a panel of nine frontier LLM judges on the bench and computed its *effective sample size* — the number of genuinely independent votes the panel behaves like, once you account for how often the judges fail together. The answer was **2.18**.

Sit with the shape of that. You budgeted for nine independent votes. You got two. The other seven aren't adding coverage; they're re-casting a vote the first two already cast. Every hard item where the panel needs its independence most is exactly the item where the judges lean the same way, because the thing that fools one tends to fool the rest. The agreement you read as *certainty* is, on those items, the same failure repeated nine times.

One object to carry through the rest of this piece: **n_eff**, the effective number of independent judges. It is not accuracy and it is not calibration. A perfectly calibrated panel can still have $n_{\text{eff}} \approx 2$ — calibration asks whether each judge is honest, **n_eff** asks whether you have as many of them as you're paying for. Kohli's 2.18 is a Kish-style effective-n computed on the judges' *error correlation* (how often they're wrong on the same items together), not on Cohen's κ and not on any accuracy gap. Keep those wires uncrossed and the rest of this follows.

The instinct is to buy different vendors. That's the wrong axis.

Everyone's first fix is the same: the judges correlate because they're too similar, so make them different — swap in GPT, Claude, Gemini, spread the panel across logos. It feels like buying independence. It is buying diversity on the axis that happens to be free of it. Think of the fix as a ladder, and notice how little the bottom rung actually gets you.

Rung 1 — different brand (GPT vs Claude vs Gemini): buys $\approx$ nothing. This is the rung everyone climbs first and it is nearly a no-op. The Frontier Paradox study (arXiv 2609.22512) measured error-correlation across judges and found that *cross-provider* correlation sits right on top of *within-provider*: cross-provider $\bar{\rho}$ lands in the ≈ 0.42 – 0.56 range, squarely overlapping within-provider (within-Gemini ≈ 0.40). Buying a different logo moved the correlation by essentially nothing. (These are measured associations, not a demonstrated cause — the study is observational, and it does not establish *why* the correlations land where they do.) In experimental terms, vendor

behaves like a *registered null*: the variable you'd expect to matter, tested, and it doesn't move the needle. A different vendor given a comparable model adds ~nothing you can bank.

Rung 2 — different architecture: buys *some* (partial, weakly supported), and it is not enough. Change the model *family*, not just the badge, and the co-failure does drop somewhat. The direct evidence for this rung is Nogueira's N-version analysis (arXiv 2607.02808): heterogeneous-model ensembles captured about **0.43–0.44** of the reliability gain true independence would deliver, versus under **0.30** for same-model ensembles. That gap — 0.43 against 0.30 — is the *delta* architecture diversity buys you: real, positive, and less than half of what full independence would give. So take this rung, but do not book it as a solution; the evidence puts it at partial, not proven, and nowhere near the confidence of rung 1.

The reason it can't reach bottom shows up when you push architecture as far as it goes. Even a *non-LLM* typed classifier — about as different an architecture as you can bolt onto a panel — does not decorrelate. Rao & Callison-Burch (arXiv 2609.29769) found a typed classifier still co-failing with LLM judges at **+45.7 percentage points over a matched-marginal null** — i.e. 45.7 points more joint failure than you'd see if their errors were independent given each judge's own failure rate. Read Nogueira and Rao–Callison-Burch together and they measure the same thing from two sides: Nogueira is the *delta* architecture diversity buys (rung 2, positive but partial); Rao–Callison-Burch is the *residual* co-failure left over *after* you have changed architecture all the way to a non-LLM (rung 3, large). Ceiling and delta, not contradiction. The correlation is not living in the architecture.

Rung 3 — the co-failure lives deeper, in shared conventions. The rungs nest for a reason: brand $\subset$ architecture $\subset$ convention, because each names a strictly *larger* set of shared causes a judge inherits — a shared brand is a subset of shared architecture, which is a subset of shared training conventions — so independence bought at a shallower rung cannot dissolve a common cause that lives at a deeper one. If the co-failure survives a change of brand *and* a change of architecture down to a non-LLM classifier, the common cause has to be something all of them share regardless of how they're built: the *scoring conventions* they learned. Annotation norms, rubric lineage, the labeling conventions baked into training data — the shared idea of "what a good answer looks like." That is the hidden common cause the ladder points at. Be precise about the epistemics here: this is a hypothesis, not a demonstrated mechanism. The evidence is *consistent with* shared conventions driving co-failure — the co-failure persists exactly where you'd expect if conventions were the culprit — but no one has run the ablation that would prove it. You could: vary the conventions, hold everything else fixed, watch the correlation move. Until someone does, treat rung 3 as the best-supported explanation on offer, not a settled fact. ("Shared conventions" is my framing of the mechanism, not a claim Rao & Callison-Burch make in those words.)

The ladder is the whole argument in one picture. Diversity is a real cure, bought on the right axis. Brand is the wrong axis. Architecture is a better axis that still doesn't reach bottom. The bottom is convention, and that's where a panel is most likely to be secretly one worldview wearing nine name tags. The rest of this piece defends each rung — the number first, then why it collapses by construction, then the axis where diversity *does* help, then a check you can run today.

The same failure shows up at pipeline scale, outside eval panels entirely. Two frontier labs — Anthropic and Meta — ran cyber-safety evaluations through the same third-party vendor, Irregular, and both failed through one shared misconfiguration. Two labs, one common cause, correlated failure — the panel problem at the level of an entire safety pipeline.

Why judges *might* correlate by construction — the mechanism behind rung 3

The reason this isn't a bug you route around by shopping vendors: the correlation is structural. This section is the mechanism story for rung 3 — where the shared-latent, shared-confounder language of representation learning turns out to be another name for the "shared conventions" the ladder points at.

Start with a corroborating observation from outside eval panels, because it isolates the "before the task started" part of the claim. In August 2026 Anthropic published "Patterns and Problems in Emerging Multiagent Systems" (Zou et al., anthropic.com/research/multiagent-systems, 2026-08-13), documenting what they call low-variance or systemic failures across agent populations. Eighteen of thirty agents, asked to name a git branch, chose the *exact same* name. Pricing agents converged to the penny — and kept converging even after the authors removed all communication between them. No messages passed. They still agreed. That last detail is the whole point: when agents converge with the channel cut, the correlation was already inside them, baked in before the task started by the thing they have in common. The report describes this vividly and never once puts a number on it — no effective-sample-size figure, no correlation coefficient. Anthropic saw the phenomenon and reached for the vocabulary ("systemic," "low variance"); the instrument to quantify it — `n_eff` — wasn't in the drawer. That instrument is exactly what this piece is about.

There's a result from representation learning worth knowing here — the Platonic Representation Hypothesis (Huh et al., arXiv 2405.07987). Its claim is that as models get larger and more capable, their internal representations *converge* — different architectures, different training runs, different labs, drifting toward the same latent geometry. I want to be precise about what that paper does and does not say, because it's easy to overreach: it's about representational convergence, and it is silent on correlated *errors*. It does not measure judge co-failure. What it gives me is a candidate mechanism for rung 3 — why judges would share a latent view of the world in the first place, regardless of brand or architecture. Read alongside the ladder, "shared latent" is the representation-learning name for "shared convention." Keep in mind this remains the rung-3 hypothesis: a plausible mechanism, not a proven one.

The step from "shared latent" to "correlated errors" is my reasoning, not theirs, so take it as reasoning: if two judges have converged onto nearly the same internal representation of "what a good answer looks like," then the inputs that fool one will tend to fool the other, because they're reading the item through the same lens. Shared representation is a shared attack surface.

There's direct evidence that the shared component is real and separable, not just a story I'm telling — and it is another rung-3 candidate. CARE (Zhao et al., arXiv 2603.00039) models judge scores as a shared quality factor Q plus latent confounders — verbosity, style, training artifacts — via a sparse-plus-low-rank precision decomposition, pulling the confounders apart from genuine quality *label-free*, and reports up to 26.8% aggregation-error reduction on UltraFeedback. That's the shared-convention idea made measurable: judges share confounders, and you can peel them off without ground truth. One honest limit — its identifiability leans on distinct eigenvalues and orthogonal loadings, so it doesn't separate *all* confounders, only the ones that geometry can pry apart. And CARE demonstrates that a shared component *exists and is separable*; it does not prove that shared conventions are *the* driver of judge co-failure. That step is still rung 3's open hypothesis.

This is what the ladder was diagnosing. If the driver of correlation were *identity* — same company,

same weights — then buying across vendors (rung 1) would fix it. But if the driver is a shared *framework* — comparable training data at web scale, convergent objectives, a common latent — then vendor is only a proxy for the mechanism, and a leaky one, which is exactly why rung 1 buys $\approx$ nothing. You can buy nine different logos and still be buying one worldview nine times.

The honest number

I said "about two." Let me show you why I trust that, and be careful about *how* I'm trusting it, because the fast way to be wrong here is to chain three different measurements as if they were one.

Three groups, three different instruments, land in the same neighborhood:

- **Kohli (arXiv 2605.29800, Apple ML):** effective sample size **2.18** for nine frontier LLM judges — Kohli's 2.18, the number I opened with. It's a Kish-style effective-n computed on the judges' error correlation.
- **Vega-Barbas (arXiv 2608.00285):** effective-diversity number (Vendi / Hill-q1) $\approx$ **1.69** for sixteen LLMs. Different estimand, different math, batch setting.
- **Noonan (arXiv 2608.21262):** And there's a related but distinct penalty at the decision boundary. Noonan's Exceedance Design Effect shows that when you care about a pass/fail cutoff on a *score* — not the vote-agreement effective-n above, a different object — the effective sample size at that tail is lower still. Different estimand, same direction: redundancy is worse exactly where a threshold decision is made.

These are not the same quantity and I'm not going to pretend they are. Kish-on-agreement, a Vendi diversity number, and Noonan's score-threshold design effect are genuinely different quantities. The point is the *convergence*: three operationalizations, built by different people for different reasons, all say your double-digit panel is behaving like a low-single-digit one. When independent instruments agree, I believe the neighborhood even when I'd argue about the street address. Call it 2, maybe a shade under, and lower at the threshold.

Three quantities nobody names cleanly

Here's a distinction I keep having to redraw in conversations, because collapsing it is how people talk past each other.

The ceiling. Given a fixed pairwise correlation ρ among your judges, adding more judges doesn't buy unbounded independence. Effective sample size saturates at $1/\rho$. Twenty judges at $\rho = 0.5$ don't give you twenty, or even ten — they give you about two, and the twenty-first buys essentially nothing. The field now has a name for this saturation: the "Correlation Ceiling" (Bay & Yearick, arXiv 2606.28661). Guard: their setting is *within-model self-consistency* on GSM8K and MATH — one model resampled — not a cross-judge panel. So I cite them for the *concept* of the ceiling, not as a measurement of panel co-failure. The math of saturation is general; their experiment is a different object.

The operating point. Where you actually sit today: $n_{\text{eff}} = n / (1 + (n-1)\rho)$. This is the standard Kish design-effect formula — it’s not owned by anyone, don’t let anyone attach a name to it. Plug in your n and your measured ρ and you get the honest count of judges you’re paying for versus receiving.

The drift monitor. Whether that operating point is *silently sliding* over time. Vendors push new checkpoints. Your rubric evolves. The mix of items you’re judging shifts. Any of these can quietly raise ρ and drop your effective- n , and a static measurement taken at deployment won’t catch it. Watching that requires a sequential test that stays valid no matter when you look — more on that below.

The field named the ceiling. It computes the operating point when asked. Nobody is monitoring the drift.

Why the vendor-diversity buy fails: the ladder pays off

Everything so far — the collapse to two, the ladder, the mechanism — comes to a head in one practical verdict: *this is why buying diversity along the vendor axis fails*. Rung 1 is the wrong axis, rung 2 is a better axis that still doesn’t reach bottom, and the reason both fall short has a forty-year-old name from software safety, not machine learning.

In 1986, Knight and Leveson (IEEE TSE) ran the experiment everyone building redundant systems should know. They had twenty-seven programs written *independently* — different programmers, no communication, the textbook conditions for independent failures — and then tested whether the programs’ failures were actually independent. They weren’t. The programs co-failed on the same inputs far more than independence predicted. The teams hadn’t colluded. They’d shared something more insidious than collusion: a *specification*. The same problem statement steered independent minds into the same corners.

That is rung 3, in 1986 clothing. The specification is the 1986 *analogue* of the shared convention — with one disanalogy worth stating plainly: Knight–Leveson’s spec was a human-written formal document handed to every team, while a shared training convention is an emergent norm distilled from web-scale data. Not the same kind of object, even if they play the same role in the argument. Your judges didn’t collude, and — as the ladder showed — they don’t even need to share a vendor or an architecture to co-fail. They share a rubric and a training objective, and those would plausibly play the role the specification played in 1986: the common cause that independent *construction* doesn’t dissolve, because it lives below construction.

Now the honest caveat, because this is exactly the place I could overclaim, and it is the same caveat that pins rung 3 to hypothesis rather than fact: I do **not** have evidence that the rubric *is* the common cause. That’s a hypothesis, an untested lever. It’s plausible and it’s instrumentable — you could vary the rubric and watch ρ move — but nobody has done that experiment cleanly, so I’m flagging it as something worth building an instrument for, not something I can assert. Knight–Leveson is an *analogue* that motivates rung 3; it is not a measurement of it. And rung 2 (Nogueira’s 0.43 vs 0.30 above) already told us architecture diversity is real but partial — so the correct reading is the ladder’s: diversify across vendors and architectures, yes, but don’t book the full independence benefit, because you’re getting less than half of it, and the rest is stuck at rung 3.

When correlation actually helps — the scoped concession

I owe you the other side, because "correlation is always bad" is false and you'd be right to distrust anyone who says it. There are two genuinely different ways the correlation stops hurting you, and it's worth being precise about both — including one framing mistake I want to head off.

Start with the case where the shared component is *noise*, not bias. If you are doing offline A/B comparison — model A versus model B, same judges scoring both — and each judge's error is *unbiased noise*, then correlation across judges is close to harmless. Noise averages out. You're taking a difference, and shared zero-mean wobble largely cancels in the difference. TensorZero has made a version of this argument and it's correct in that scope. If that's your entire use case, you can file this article under "interesting, not urgent."

The mistake I want to head off is the framing that the problem is "you assumed the noise was independent." That's not the problem. The problem is that you never modeled the *joint* distribution of the judges' verdicts on the *same item* at all. Independence was never a modeling assumption you made and could relax — it was a silent default in the arithmetic of averaging. And it stops being harmless the moment the shared component is *bias*, not noise, and the moment you're doing live gatekeeping rather than a symmetric comparison. On a hard item where the training distribution nudges every judge the same wrong way, averaging nine correlated wrong answers gives you a very confident wrong answer. The panel's agreement reads as certainty. It's actually just shared blind spots, voting. So the thesis is scoped: systematic shared bias, on hard items, in a gatekeeping decision. That's the regime. Inside it, the deficit is real.

The second case is the voting regime. Kaniovski and Zaigraev (Theory and Decision 71:439–459, 2011) worked out the jury case exactly. Above a super-majority quota, *positive* correlation among jurors can *increase* the panel's verdict accuracy. The sign of the effect flips with the decision rule. So there are voting regimes where your judges agreeing is a feature.

Two things keep this from being an endorsement of correlated panels. First, even in the best case they analyze — unanimity — a single competent juror beats the whole correlated panel. So we're talking about the sign of a marginal effect, not a reason to prefer the crowd. Second, and this is the operational sting: the benefit shows up under strict quotas, and the deficit is *worst* under simple majority vote — which is exactly what practitioners actually deploy. The regime where correlation helps is not the regime you're running.

And if the shared signal is reliably *right* — this is the Ross/Lo point (arXiv 2609.04373) — then judges agreeing is efficiency, not fragility. Conceded, fully. That's why the thesis is scoped to the miscalibrated-on-hard-items regime. When the panel is right, its agreement is free money. The problem is that agreement looks identical whether the panel is right or collectively fooled, and you can't tell which from the agreement alone.

Two objections I have to meet head-on

"More agents is all you need." There's a well-known result (arXiv 2402.05120) that resampling *one* model many times and voting gives you real gains. True — and not a counterexample. Resampling one model draws near-i.i.d. samples *by construction*; the independence is manufactured by the sampling process. Cross-*model* panels are the opposite case: they don't get that independence, precisely because the models share the training specification the resampling trick doesn't have to

worry about. Same voting math, opposite independence structure. The i.i.d. case is the one where the panel really is nine; the cross-model case is the one where it's two.

"The market will differentiate the judges for you." Competition should push producers to differentiate — Raghavan (arXiv 2412.08610) is the reference for that dynamic. But a judge panel is a *single-buyer* market. One team picks all nine judges. The differentiation incentive fires when many buyers pull producers in many directions; when a single buyer assembles the panel, nothing forces the members apart, and the shared-specification pull wins by default. The market can't decorrelate a panel one person shopped for.

The instrument that fills the hole

So how do you measure co-failure when you have no ground-truth labels — which, live, you almost never do? This is the monitor I've been building, and I'll keep it honest by keeping it high-level: it's research-stage, not a shipped product.

The core trick is to pair judges' verdicts *across different items*. Two different items are independent by construction — whatever the panel's shared bias is, it can't correlate a verdict on item 5 with a verdict on item 200 the way it correlates two verdicts on item 5. So cross-item pairing gives you a *label-free, margins-free* baseline: the co-agreement you'd expect if the judges were independent, estimated without ever needing to know the right answer. Then you compare the same-item agreement against that baseline. Excess same-item agreement over the cross-item baseline is your co-failure signal.

This also sidesteps a confound that sinks the naive approach. Jo, Garg & Raghavan ("The Subjectivity of Monoculture", arXiv 2602.24086) show that item-exchangeable independence nulls — the $p_1 \cdot p_2$ product baseline — overstate correlation, because item *difficulty* induces agreement on its own: everyone gets the easy ones right and the hard ones wrong, and that shared response to difficulty looks like co-failure even under genuine independence. Cross-item pairing dodges this by construction: pairing verdicts on *different* items, which are independent whatever their difficulties, rather than trying to correct the null after the fact. Note their critique targets the $p_1 p_2$ null (Kim/Goel), not the Kish-ICC n_{eff} — so it's complementary to the effective-sample-size framing here, not a threat to it.

One competitor I have to name, because I'm not the only one measuring judge correlation without labels. Sunkavalli (arXiv 2609.08826) gives a closed-form, *reference-free* estimator of judge error-correlation under a single-common-factor model. It's real prior art and I'll name the differences precisely, because the differences are the whole moat: (a) Sunkavalli needs a *second anchor-evaluator panel* — at least two judges *and* at least two anchor evaluators — where the approach here works from a single judge panel's own outputs; (b) it's batch, one-shot, where this is anytime-valid and sequential; (c) it pools covariances within the item set, with no cross-item pairing. Different information requirement, different validity regime, different estimator. But same honest goal, and worth knowing about.

Run the cross-item test as an *anytime-valid sequential test* — an e-process — and you get the drift monitor from three sections ago. Anytime-valid means you can watch the statistic continuously and stop or alarm whenever it crosses, without inflating your error rate by peeking. No fixed sample size, no p-hacking your stopping time. That's the property you need for a monitor that runs forever against a stream of judgments and a vendor that ships a new checkpoint on a Tuesday.

The instrument is research-stage; a production-hardened implementation doesn't exist yet. But the shape is what I'm confident about, and the shape is enough to build on.

One scope stake in the ground: this monitor assumes *non-interacting parallel panels* — judges scoring independently, votes pooled after. If your architecture is a *debate* — judges arguing across rounds, seeing each other's reasoning — you're in a different and worse regime, and this instrument's independence baseline doesn't apply. Ma et al. ("Judging with Many Minds", arXiv 2505.19477) show that debate sharply amplifies bias after the *first* round and sustains it thereafter — across position, verbosity, bandwagon, and chain-of-thought biases — while parallel meta-judge pools stay comparable to a single judge. I want to be careful about what that does and doesn't license: they compute no effective-sample-size and no correlation coefficient, so I can't tell you debate drives n_{eff} below anything. What they show is directional and it's enough for a design rule: interaction amplifies the shared bias rather than diluting it. Keep the panel parallel.

Compute it yourself

I've been telling you to measure a number. Here's how, with no dependency on my research-stage monitor — just your judge-score matrix and fifteen lines of numpy. There are two formulas, and naming both up front kills the "which one?" objection, because they answer different questions at different times.

ESDOF / participation ratio — the design-time screen. Before you commit to a panel, you want to know how many independent opinions its correlation structure can *possibly* deliver. Take the judge $\times$ judge correlation matrix, get its eigenvalues, and compute

$$n_{\text{eff}} = (\sum \lambda_i)^2 / \sum \lambda_i^2$$

the ratio of the squared sum of eigenvalues to the sum of squared eigenvalues. This is the Hill $q=2$ number. It's not new — it's the effective-degrees-of-freedom formula from Bretherton et al. 1999 (*J. Climate*), operational at ECMWF for decades. It needs only the correlation matrix, which makes it the right pre-hoc screen: you can run it on a pilot before you've collected per-item labels.

Kish — the validation-time check. Once the panel is live and you have per-item judge agreement data, compute the mean pairwise co-failure correlation $\bar{\varphi}$ across your items (this $\bar{\varphi}$ is the same ρ from the design-effect formula above — the mean off-diagonal pairwise correlation) and plug into the design-effect formula from the last section:

$$n_{\text{eff}} = n / (1 + (n-1) \cdot \bar{\varphi})$$

This is post-hoc panel validation: it uses the disagreement you actually observed rather than the correlation matrix's eigenstructure. ESDOF screens a design; Kish audits a deployment.

Both, from a judge-score matrix S of shape (items $\times$ judges):

```
import numpy as np
```

```

def n_eff_esdof(S):
    # design-time screen: participation ratio of the judge-judge correlation matrix
    C = np.corrcoef(S, rowvar=False)          # judges x judges
    eig = np.linalg.eigvalsh(C)
    eig = np.clip(eig, 0, None)                # guard tiny negative eigenvalues
    return eig.sum()**2 / np.square(eig).sum()

def n_eff_kish(S):
    # validation-time: Kish design effect from mean pairwise correlation
    n = S.shape[1]
    C = np.corrcoef(S, rowvar=False)
    iu = np.triu_indices(n, k=1)
    phi_bar = C[iu].mean()                    # mean off-diagonal (pairwise)
    correlation
    return n / (1 + (n - 1) * phi_bar)

# S[i, j] = judge j's score on item i
print("ESDOF:", n_eff_esdof(S), " Kish:", n_eff_kish(S))

```

As far as I can find, no open-source LLM-judge `n_eff` audit tool exists — the `judges` PyPI package, JudgeLM, and OpenCompass all measure per-judge quality and none of them measure panel independence. The closest precedent is ASSERT-KTH's Knight-Leveson-Redux (github.com/ASSERT-KTH/Knight-Leveson-Redux), which re-runs the 1986 N-version experiment on AI *coding* agents: a million test cases, per-version failure rate, a coincident-failure z-statistic, pairwise φ , a Bonferroni binomial test. Its domain is code generation, not judging, but the methodology transfers directly — swap "program fails a test case" for "judge misgrades an item" and it's the same audit.

The one column to add

Every practitioner guide I've read asks: is each judge calibrated? That's the κ question, and it's a good question. None of them ask: are your nine judges actually nine? That's the effective-sample-size question, and it's the one with your money in it.

So add the column. Run the Kish check from "Compute it yourself" on your panel and put `n_eff` on the same dashboard as your κ . One number. It will, I'd wager, embarrass your headcount.

One measurement note while you're at it. For binary verdicts, $\kappa = q \cdot \varphi$ (Rao & Callison-Burch, arXiv 2606.00093), where φ is the raw agreement correlation and q depends on your label distribution. When your pass/fail split is skewed — and gatekeeping splits usually are, most things pass — that q factor drags κ down, so κ *understates* the agreement structure you actually have. Report φ , not just κ , or you'll flatter yourself about how independent the panel is.

The thing to sit with: your LLM passes the tests and can still fail the audit. Calibration and independence are different properties, and every guide in the field checks the first and none check the second. You've been auditing whether each judge is honest. You haven't been auditing whether you have as many of them as you're paying for. Add the column.